

Random Forests

Günter J. Hitsch

February 21, 2017

```
library(data.table)
library(ggplot2)
library(ranger)
```

Data

As in the regression tree example, we use household-level private label share data aggregated at the *year* level.

```
load("./Data/PL_shares_annual.RData")
```

For easier interpretability, we convert income and the Zillow housing index to logs.

```
PL_shares_DT[, `:=`(income      = log(income),
                    zillow_index = log(zillow_index))]
setnames(PL_shares_DT, c("income", "zillow_index"), c("log_income", "log_zillow"))
```

Using ranger

The **ranger** function will throw an error if it detects missing data values(NA). Hence, we first remove observations with missing values:

```
PL_shares_DT = PL_shares_DT[complete.cases(PL_shares_DT)]
```

Now estimate a random forest:

```
fit = ranger(PL_share ~ . - log_zillow,
             data = PL_shares_DT[, !"household_code", with = FALSE],
             num.trees = 2000,
             # importance = "impurity",
             seed = 1776)
```

Options:

- **num.trees** specifies the number of trees to incorporate. When you first run the algorithm, use a small number to get a sense of how slow or fast the trees are grown (**ranger** will provide some corresponding output, unless you add the option **verbose = FALSE**). The default setting is **num.trees = 500**. In a final production run you may choose a larger number, maybe 1000 or more if time permits.
- Set a **seed** inside ranger to exactly replicate your previous results.
- The **importance = "impurity"** option allows you to later use the **importance** function (e.g. **importance(fit)**) to display the variable importance measures.

Note

Don't exclude variables inside a formula when using **ranger**! E.g., do not use

```
PL_share ~ . - household_code
```

Instead, exclude the variables directly from the data.table:

```
DT[, !c("var_1", "var_2"), with = FALSE]
```

Random forest prediction

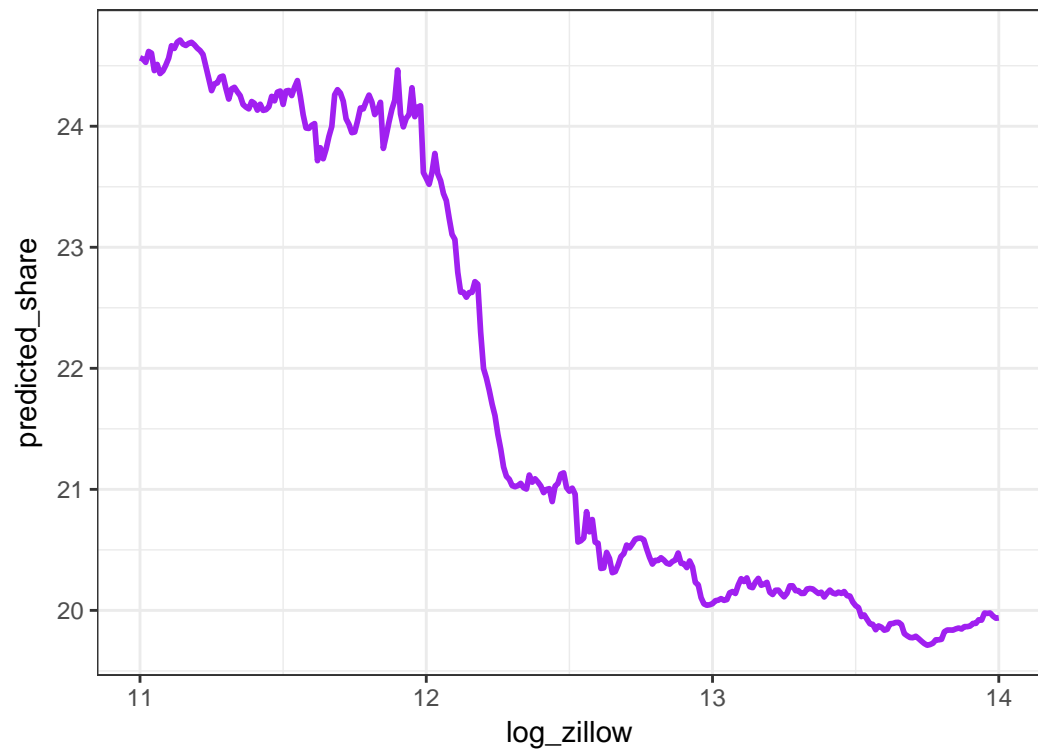
Create a `data.table` with a range of Zillow index values.

```
predict_DT = data.table(log_zillow = seq(11, 14, 0.01),
  log_income = 9.5,
  year = 2012,
  unemployed = FALSE,
  education = "Graduated High School",
  age = 55,
  size = 2,
  has_children = 1,
  female_head = FALSE,
  marital_status = "Married",
  race = "White",
  hispanic_origin = "Hispanic")
```

Predict and graph the predicted private-label share vs. the Zillow index:

```
predict_ranger = predict(fit, data = predict_DT)
predict_DT[, predicted_share := predict_ranger$predictions]

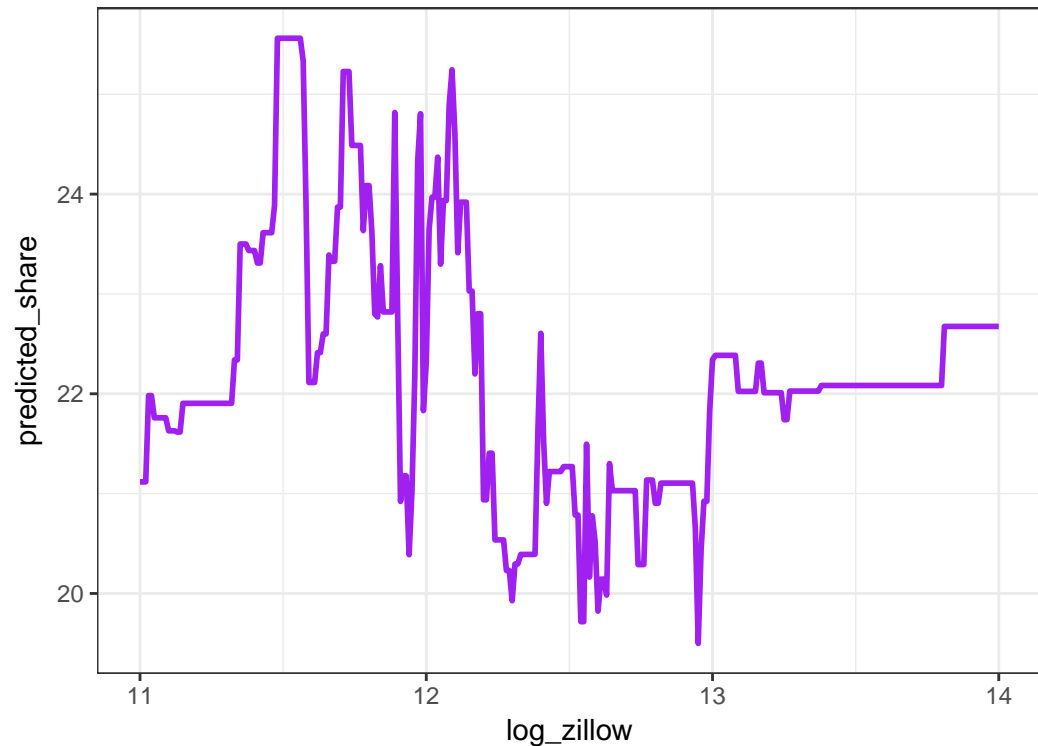
ggplot(predict_DT, aes(x = log_zillow, y = predicted_share)) +
  geom_line(color = "purple", size = 1) +
  theme_bw()
```



Here we use only 25 trees in the prediction:

```
predict_ranger = predict(fit, data = predict_DT, num.trees = 25)
predict_DT[, predicted_share := predict_ranger$predictions]
```

```
ggplot(predict_DT, aes(x = log_zillow, y = predicted_share)) +
  geom_line(color = "purple", size = 1) +
  theme_bw()
```



pdf graphs

Predict a matrix with columns for each individual tree using the `predict.all = TRUE` option.

```
predict_ranger = predict(fit, data = predict_DT, predict.all = TRUE)
```

First, plot each individual tree.

```
L = 100
```

```
pdf(file = "Individual-Trees.pdf", width = 8, height = 6)
for (i in 1:L) {
  predict_DT[, predicted_share := predict_ranger$predictions[, i]]

  plot_i = ggplot(predict_DT, aes(x = log_zillow, y = predicted_share)) +
    geom_line(color = "purple", size = 1) +
    scale_y_continuous(paste0("PL share, tree: ", i)) +
    theme_bw()
  print(plot_i)
}
dev.off()
```

Second, predict based on the average over the first `i` trees:

```

L = 2000

pdf(file = "Tree-Averages.pdf", width = 8, height = 6)
i = 1
while (i <= L) {
  if (i < 2) predict_DT[, predicted_share := predict_ranger$predictions[, 1]]
  else predict_DT[, predicted_share := rowMeans(predict_ranger$predictions[, 1:i])]

  plot_i = ggplot(predict_DT, aes(x = log_zillow, y = predicted_share)) +
    geom_line(color = "purple", size = 1) +
    scale_y_continuous(paste0("PL share, trees: ", i)) +
    theme_bw()
  print(plot_i)

  cat(i, "\n")
  i = i + 1 + floor(sqrt(i)/5)
}
dev.off()

```